

CS/ECE/ME 532

Homework 7: Convexity and the SVM

1. Verifying convexity. Prove that the following functions are convex.

- a) The sum of two convex functions: $f(\mathbf{x}) = g(\mathbf{x}) + h(\mathbf{x})$ where g and h are convex.
- b) A positive quadratic form: $f(\mathbf{x}) = \mathbf{x}^\top \mathbf{P} \mathbf{x}$, where $\mathbf{P} \succ 0$.
- c) The pointwise maximum of several affine functions: $f(\mathbf{x}) = \max_{i \in \{1, \dots, m\}} (\mathbf{a}_i^\top \mathbf{x} + b_i)$
- d) The definition of convexity we saw in class may also be extended to functions that take a matrix as an argument, e.g. $f : \mathbb{R}^{m \times n} \rightarrow \mathbb{R}$. Prove that $f(\mathbf{X}) = \|\mathbf{X}\|_2$ (the induced 2-norm) is convex.

2. Gradient Descent and Stochastic Gradient Descent. Suppose we have training data $\{\mathbf{x}_i, y_i\}_{i=1}^m$, with $\mathbf{x}_i \in \mathbb{R}^n$ and y_i is a scalar label. Derive gradient descent and SGD algorithms to solve the following ℓ_1 -loss optimization:

$$\min_{\mathbf{w} \in \mathbb{R}^n} \sum_{i=1}^m |y_i - \mathbf{x}_i^\top \mathbf{w}|.$$


- a) Simulate this problem as follows. Generate each x_i as random points in the interval $[0, 1]$ and generate $y_i = w_1 x_i + w_2 + \epsilon_i$, where w_1 and w_2 are the slope and intercept of a line (of your choice) and $\epsilon_i = \text{randn}$, a Gaussian random error generated in Matlab. With $m = 10$. Repeat this experiment with several different datasets (with different random errors in each case).
- b) Implement the GD or SGD algorithm for the ℓ_1 -loss optimization. Compare the solution to this optimization with the LS line fit.
- c) Now change the simulation as follows. Instead of generating ϵ_i as Gaussian, now generate the errors according to a Laplacian (two-sided exponential distribution) using `laprnd(1,1)`. Compare the LS and ℓ_1 -loss solution compare in this case. Repeat this experiment with several different datasets (with different random errors in each case).

3. Error Bounds using Hinge Loss. State the SGD algorithm for solving the hinge-loss optimization

$$\min_{\mathbf{w}} \sum_{i=1}^m f_i(\mathbf{w}) \quad \text{where:} \quad f_i(\mathbf{w}) = (1 - y_i \mathbf{x}_i^\top \mathbf{w})_+.$$

- a) Derive a bound on the average error $\frac{1}{T} \sum_{t=1}^T (f_{i_t}(\mathbf{w}_t) - f_{i_t}(\mathbf{w}^*))$ using Theorem 1 from the lecture notes (on moodle). Assume that $\mathbf{w}_1 = \mathbf{0}$ and $\|\mathbf{w}^*\| \leq 1$, and that the features are normalized so that $\|\mathbf{x}_i\| \leq 1$ for all i . Assume a constant stepsize of $\gamma = 1/\sqrt{T}$ as in Corollary 1.
- b) How many iterations are required to guarantee that the average error is less than 0.01?

4. **Classification and the SVM.** Revisit the iris data set from Homework 3. For this problem, we will use the 3rd and 4th features to classify whether an iris is *versicolor* or *virginica*. Here is a plot of the data set for this restricted set of features.



We will look for a linear classifier of the form: $x_{i3}w_1 + x_{i4}w_2 + w_3 \approx y_i$. Here, x_{ij} is the measurement of the j^{th} feature of the i^{th} iris, and w_1, w_2, w_3 are the weights we would like to find. The y_i are the labels; e.g. +1 for *versicolor* and -1 for *virginica*.

- Reproduce the plot above, and also plot the decision boundary for the least squares classifier.
- This time, we will use a regularized SVM classifier with the following loss function:

$$\underset{\mathbf{w}}{\text{minimize}} \quad \sum_{i=1}^m (1 - y_i \mathbf{x}_i^T \mathbf{w})_+ + \lambda(w_1^2 + w_2^2)$$

Here, we are using the standard hinge loss, but with an ℓ_2 regularization that penalizes only w_1 and w_2 (we do not penalize the offset term w_3). Solve the problem by implementing gradient descent of the form $\mathbf{w}_{t+1} = \mathbf{w}_t - \gamma \nabla f(\mathbf{w}_t)$. For your numerical simulation, use parameters $\lambda = 0.1$, $\gamma = 0.003$, $\mathbf{w}_0 = \mathbf{0}$ and $T = 20,000$ iterations. Plot the decision boundary for this SVM classifier. How does it compare to the least squares classifier?

- Let's take a closer look at the convergence properties of \mathbf{w}_t . Plot the three components of \mathbf{w}_t on the same axes, as a function of the iteration number t . Do the three curves each appear to be converging? Now produce the same plots with a larger stepsize ($\gamma = 0.01$) and a smaller stepsize ($\gamma = 0.0001$). What do you observe?